




# EECS 442 – Computer vision

## Segmentation & Clustering

- Segmentation in human vision
- K-mean clustering
- Mean-shift
- Graph-cut

Reading: Chapters 14 [FP]

Some slides of this lectures are courtesy of prof F. Li, prof S. Lazebnik, and various other lecturers



# Segmentation

- Compact representation for image data in terms of a set of **components**
- Components share “common” **visual properties**
- Properties can be defined at **different level of abstractions**


# General ideas

- **Tokens**
    - whatever we need to group (pixels, points, surface elements, etc., etc.)
  - **Bottom up segmentation**
    - tokens belong together because they are locally coherent
  - **Top down segmentation**
    - tokens belong together because they lie on the same object
- > These two are not mutually exclusive

# What is Segmentation?

- Clustering image elements that “belong together”
  - Partitioning
    - Divide into regions/sequences with coherent internal properties
  - Grouping
    - Identify sets of coherent tokens in image

# What is Segmentation?




Why do these tokens belong together?

# Basic ideas of grouping in human vision

- Figure-ground discrimination
- Gestalt properties


# Figure-ground discrimination



- Grouping can be seen in terms of allocating some elements to a figure, some to ground
- Can be based on local bottom-up cues or high level recognition



# Figure-ground discrimination



# Gestalt properties

- A series of factors affect whether elements should be grouped together



Not grouped



Proximity




Similarity



Similarity




Common Fate




Common Region


# Gestalt properties




Parallelism



Symmetry




Continuity



Closure

# Gestalt properties



# Gestalt properties


Grouping  
by occlusions



# Gestalt properties

Grouping  
by invisible  
completions

A



B



C



D



# Emergence



# Segmentation in computer vision

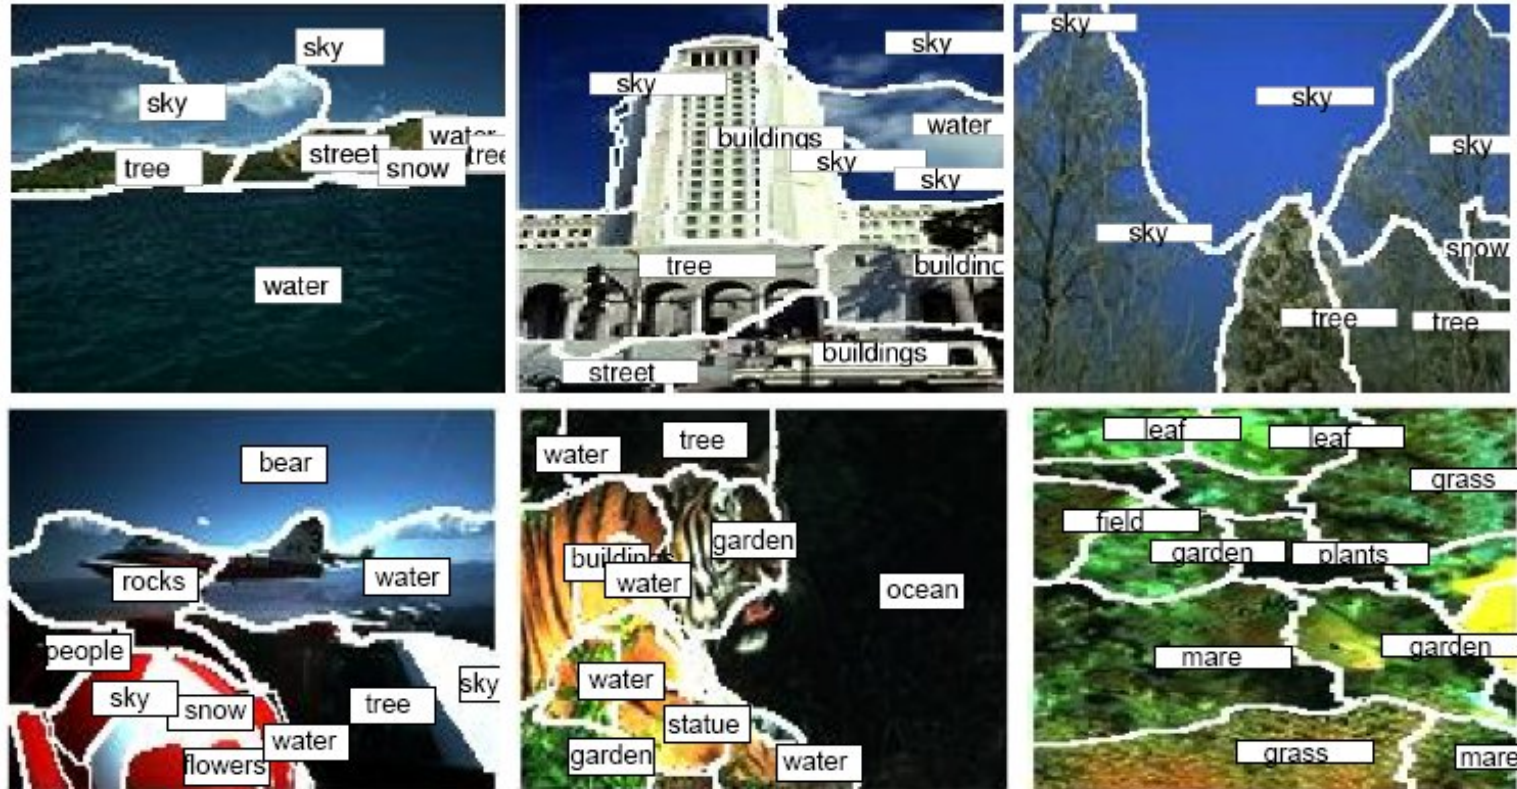
J. Malik, S. Belongie, T. Leung and J. Shi. "*Contour and Texture Analysis for Image Segmentation*". IJCV 43(1),7-27,2001.





# Segmentation in computer vision

Object Recognition as Machine Translation, Duygulu, Barnard, de Freitas, Forsyth, ECCV02



# Segmentation as clustering

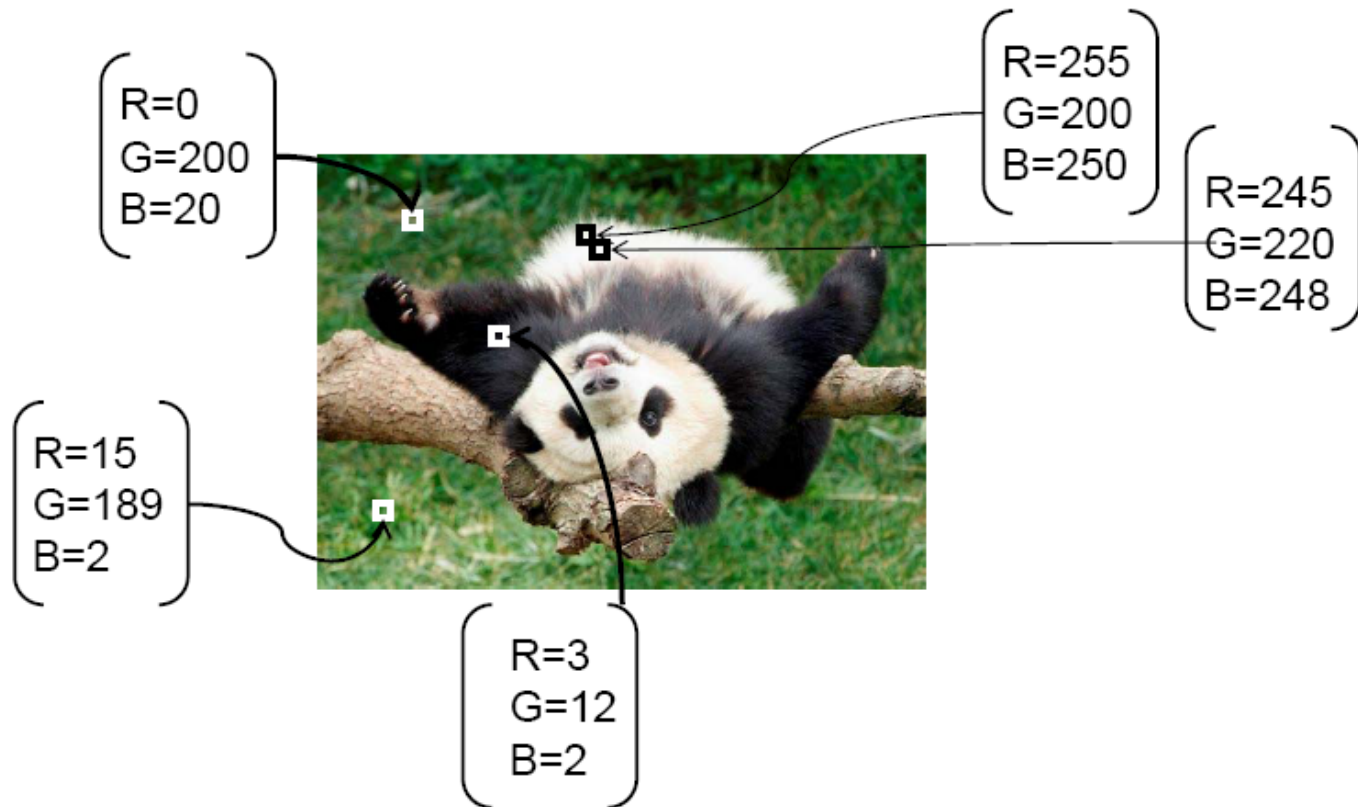
Cluster together tokens that share similar visual characteristics

- K-mean
- Mean-shift
- Graph-cut


# Feature Space

- Every token is identified by a set of salient visual characteristics. For example:
  - Position
  - Color
  - Texture
  - Motion vector
  - Size, orientation (if token is larger than a pixel)


# Feature Space



*Feature space:*  
each token is represented by a point




Token **similarity** is thus measured by distance between points (“feature vectors”) in feature space



$$\sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + \cdots + (p_n - q_n)^2} = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$$


# Cluster together tokens with high similarity






























- Agglomerative clustering
  - Add token to cluster if token is similar enough to element of clusters
  - Repeat
- Divisive clustering
  - Split cluster into subclusters if along best boundary
  - Boundary separates subclusters based on similarity
  - Repeat

# Hierarchical structure of clusters (Dendrograms)



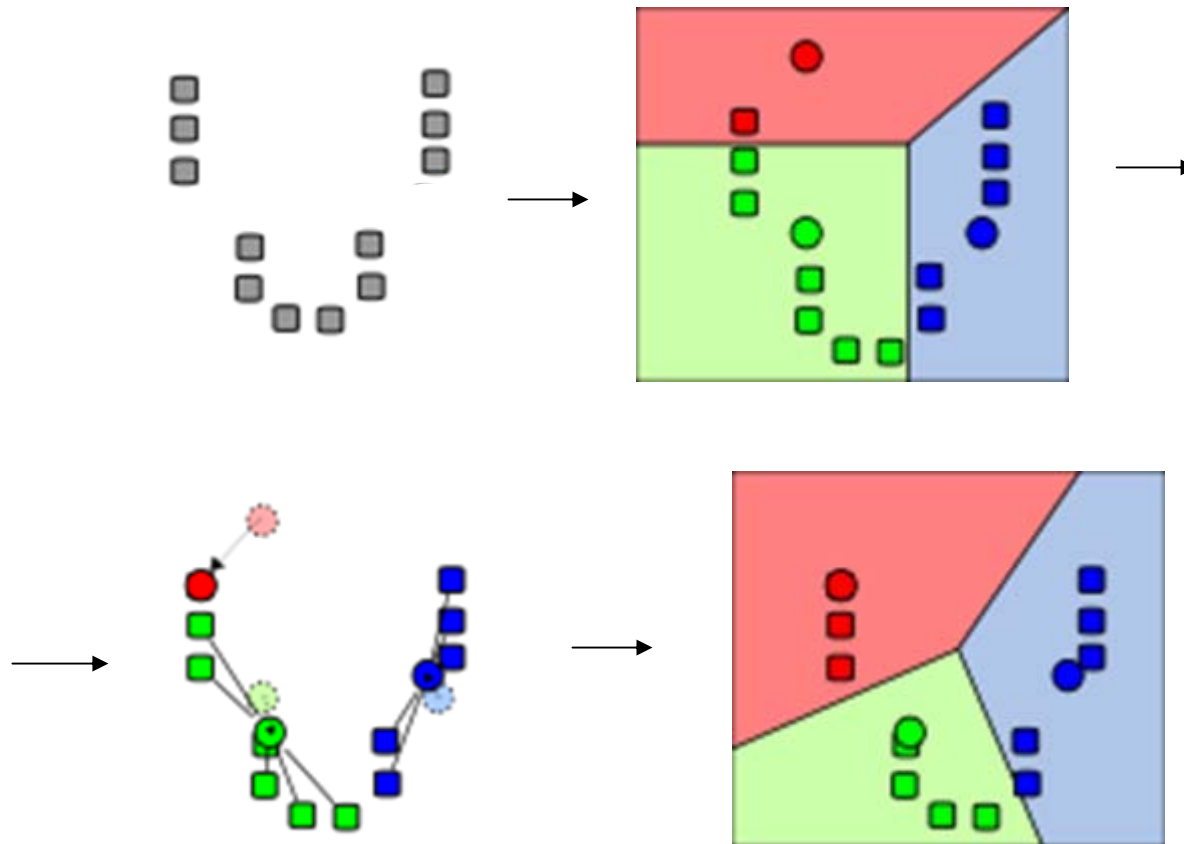
- `rand('seed',12);`
- `X = rand(100,2);`
- `Y = pdist(X,'euclidean');`
- `Z = linkage(Y,'single');`
- `[H, T] = dendrogram(Z);`



# K-Means Clustering

- Initialization: Given  $K$  categories,  $N$  points in feature space. Pick  $K$  points randomly; these are initial cluster centers (means)  $m_1, \dots, m_K$ . Repeat the following:
  1. Assign each of the  $N$  points,  $x_i$ , to clusters by nearest  $m_i$
  2. Re-compute mean  $m_i$  of each cluster from its member points
  3. If no mean has changed more than some  $\varepsilon$ , stop

# Example: 3-means Clustering





- Effectively carries out gradient descent to minimize:

$$e(\mathbf{m}_i) = \sum_{i=1}^{n_c} \sum_{j; c_j=i} |\mathbf{x}_j - \mathbf{m}_i|^2$$

$$\frac{\partial e}{\partial \mathbf{m}_k} = \sum_{j; c_j=k} -2(\mathbf{x}_j - \mathbf{m}_k) = 0$$

$$\mathbf{m}_k = \frac{\sum_{j; c_j=k} \mathbf{x}_j}{\sum_{j; c_j=k} 1} = \frac{1}{n_k} \sum_{j; c_j=k} \mathbf{x}_j$$









# K-means clustering using intensity alone and color alone

Image



Clusters on intensity



Clusters on color




# K-Means pros and cons

- Pros


- Simple and fast
- Converges to a local minimum of the error function

- Cons

- Need to pick K
- Sensitive to initialization
- Only finds “spherical” clusters
- Sensitive to outliers



(A): Undesirable clusters



(B): Ideal clusters



# Mean shift segmentation

D. Comaniciu and P. Meer, Mean Shift: A Robust Approach toward Feature Space Analysis, PAMI 2002.

- An advanced and versatile technique for clustering-based segmentation

**Segmented "landscape 1"**





**Segmented "landscape 2"**



# Mean shift segmentation

- The mean shift algorithm seeks a *mode* or local maximum of density of a given distribution
  - Choose a search window (width and location)
  - Compute the mean of the data in the search window
  - Center the search window at the new mean location
  - Repeat until convergence






# Mean shift

Region of  
interest

Center of  
mass

Mean Shift  
vector





# Mean shift

Region of  
interest


Center of  
mass

Mean Shift  
vector











# Computing The Mean Shift

Simple Mean Shift procedure:

- Compute mean shift vector
- Translate the Kernel window by  $\mathbf{m}(\mathbf{x})$



# Real Modality Analysis




- Tessellate the space with windows


- Merge windows that end up near the same "peak" or model

# Attraction basin


- **Attraction basin:** the region for which all trajectories lead to the same mode
- **Cluster:** all data points in the attraction basin of a mode




# Attraction basin



(a)





(b)




# Segmentation by Mean Shift

- Find features (color, gradients, texture, etc)
- Initialize windows at individual pixel locations
- Perform mean shift for each window until convergence
- Merge windows that end up near the same “peak” or mode





# Mean shift segmentation results



<http://www.caip.rutgers.edu/~comanici/MSPAMI/msPamiResults.html>







# Mean shift pros and cons

- Pros
  - Does not assume spherical clusters
  - Just a single parameter (window size)
  - Finds variable number of modes
  - Robust to outliers
- Cons
  - Output depends on window size
  - Computationally expensive
  - Does not scale well with dimension of feature space

# Graph-based segmentation

- Represent features and their relationships using a graph
- Cut the graph to get subgraphs with strong interior links and weaker exterior links

# Images as graphs



- Node for every pixel
- Edge between every pair of pixels
- Each edge is weighted by the *affinity* or similarity of the two nodes

# Measuring Affinity

Distance

$$aff(x, y) = \exp \left\{ - \left( \frac{1}{2\sigma_d^2} \right) (\|x - y\|^2) \right\}$$


Intensity

$$aff(x, y) = \exp \left\{ - \left( \frac{1}{2\sigma_i^2} \right) (\|I(x) - I(y)\|^2) \right\}$$

Color


$$aff(x, y) = \exp \left\{ - \left( \frac{1}{2\sigma_c^2} \right) (\|c(x) - c(y)\|^2) \right\}$$

# Segmentation by graph partitioning



- Break Graph into sub-graphs
  - Break links (**cutting**) that have low affinity
    - similar pixels should be in the same sub-graphs
    - dissimilar pixels should be in different sub-graphs
- Sub-graphs represents different image segments


# Graph cut



- **CUT:** Set of edges whose removal makes a graph disconnected
- Cost of a cut: sum of weights of cut edges
- A graph cut gives us a segmentation
  - What is a “good” graph cut and how do we find one?




# Graph cut



- CUT: Set of edges whose removal makes a graph disconnected
- Cost of a cut: sum of weights of cut edges
- A graph cut gives us a segmentation
  - What is a “good” graph cut and how do we find one?

# Minimum cut

- We can do segmentation by finding the *minimum cut* in a graph
  - Efficient algorithms exist for doing this
- Drawback: minimum cut tends to cut off very small, isolated components



# Normalized cut

J. Shi and J. Malik. Normalized cuts and image segmentation. PAMI 2000

- IDEA: normalizing the cut by component size
- The *normalized cut* cost is:

$$\frac{cut(A, B)}{assoc(A, V)} + \frac{cut(B, A)}{assoc(B, V)}$$

$assoc(A, V)$  = sum of weights of all edges in  $V$  that touch  $A$

- The exact solution is NP-hard but an approximation can be computed by solving a *generalized eigenvalue* problem

# Normalized cuts: Pro and con

- Pros
  - Generic framework, can be used with many different features and affinity formulations
- Cons
  - High storage requirement and time complexity
  - Bias towards partitioning into equal segments

# Normalized cuts: Results




Figure from “Image and video segmentation: the normalised cut framework”, by Shi and Malik, copyright IEEE, 1998

# Normalized cuts: Results




Figure 1





Figure from “Normalized cuts and image segmentation,” Shi and Malik, copyright IEEE, 2000




# Contour and Texture Analysis for Image Segmentation

J. Malik, S. Belongie, T. Leung and J. Shi. "*Contour and Texture Analysis for Image Segmentation*". IJCV 43(1),7-27,2001.



# Efficient Graph-Based Image Segmentation

Efficient Graph-Based Image Segmentation Pedro F. Felzenszwalb and Daniel P. Huttenlocher  
International Journal of Computer Vision, Volume 59, Number 2, September 2004




# Integrating top-down and bottom-up segmentation

Z.W. Tu, X.R. Chen, A.L. Yuille, and S.C. Zhu. Image parsing: unifying segmentation, detection and recognition. IJCV 63(2), 113-140, 2005.




a. An example image



b. Generic regions



c. Text



d. Faces



# EECS 442 – Computer vision

## Object Recognition

