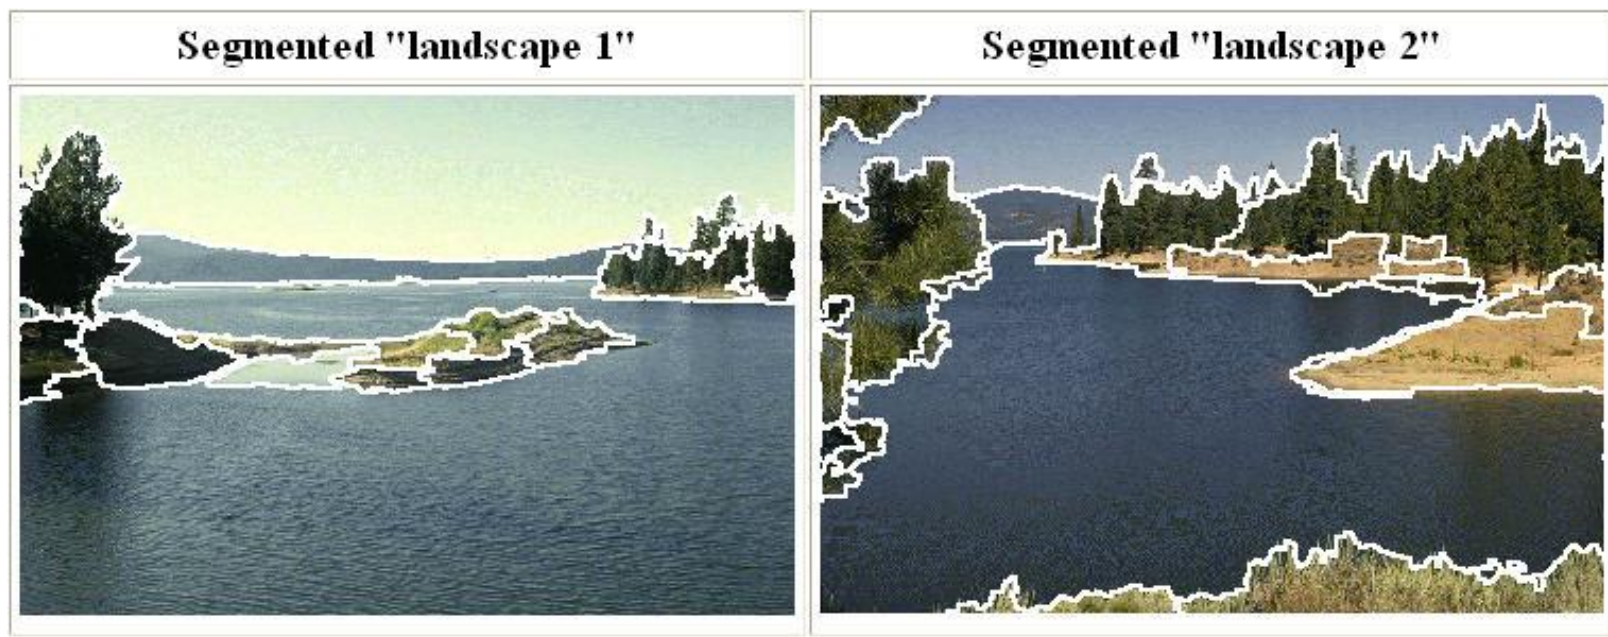


Mean shift segmentation

D. Comaniciu and P. Meer, Mean Shift: A Robust Approach toward Feature Space Analysis,
PAMI 2002, v.24, 603-619.





Versatile technique for clustering-based segmentation.

Uses the $L^*u^*v^*$ color space which is also perceptually uniform.
A nonlinear transformation from RGB.



Mean shift algorithm...

...try to find *modes* of this non-parametric density.



Kernel density estimation

Kernel density estimation function

$$\hat{f}_h(x) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right)$$

$K(x) > 0$ only for $\|x\| \leq 1$

the bandwidth, h , has to be given by the user.

The kernel is symmetric and depends on x^2 .


The Epanechnikov kernel $\sim (1 - \|x\|^2)$

and the truncated Gaussian kernel


$$K\left(\frac{x - x_i}{h}\right) = \frac{1}{\sqrt{2\pi}} e^{-\frac{(x - x_i)^2}{2h^2}}$$

are used.


Mean shift




Mean shift





Mean shift



Mean shift



Two synthetic Gaussian modes, but is not used in mean shift.
The only parameter used is h .



The two kernels converge to different modes in spite starting from overlapping regions.

Computing the Mean Shift


$k(x^2) = K(x)$ profile of the kernel

$g(x) = -k'(x)$


$$\hat{f}_h(x) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right)$$

gradient $f(x) = 0$
in each iteration

- Translate the kernel window by $\mathbf{m}(\mathbf{x})$




Modality Analysis




- *Tessellate* the space with windows.
- Merge windows that end up *near* the same mode (peak).

Attraction basin


- \mathcal{A}_i : the region for which all trajectories lead to the same mode
- $\#$ all data points in the attraction basin of a mode




Example: attraction basins



(a)



(b)



zero gradient, $g(x)=0$
but not a maximum
stationary point

eliminated
by shifting a little bit
the trajectory

A color pixel is represented by a *spatial* and a *range* bandwidth.

a two-dimensional spatial bandwidth: h_s

a three-dimensional range bandwidth: h_r

The user have to give only this two parameters. They are not very strict like in k-means.

$$K_{h_s, h_r}(\mathbf{x}) = \frac{C}{h_s^2 h_r^3} k\left(\left\|\frac{\mathbf{x}^s}{h_s}\right\|^2\right) k\left(\left\|\frac{\mathbf{x}^r}{h_r}\right\|^2\right)$$

2 + 3 = 5 dimensions

Gray level images have one-dimensional range bandwidth.

Example of a window's convergence.

256 x 256



(a)




(b)

Cameraman image. (a) Original. (b) Mean shift filtered $(h_s, h_r) = (8, 4)$.


(8,4)

original
(inverted)




(a)

mean shift




(b)

filtering



(c)

segmentation



(d)

Fig. 4. Visualization of mean shift-based filtering and segmentation for gray-level data. (a) Input. (b) Mean shift paths for the pixels on the plateau and on the line. The black dots are the points of convergence. (c) Filtering result (h_s, h_r) = (8, 4). (d) Segmentation result.

Mean shift clustering

- The mean shift algorithm seeks modes of the given set of points.
 1. Choose kernel and two bandwidths.
 2. For each point:
 - a) Center a window on that point.
 - b) Compute the mean of the data in the search window.
 - c) Center the search window at the new mean location.
 - d) Repeat (b,c) until convergence.
 3. Assign points that lead to nearby modes to the same cluster.

In segmentation, the means are both *in spatial and range* and the points always converge *to the nearest mode*.

Mean shift clustering are *two* phases:

filtering, as was described before;


segmentation, unify adjacent clusters if they are closer than h_s in the spatial domain and h_r in the range domain. (Step 3.)

EDISON program can do additional things too, but we will not describe them here.


Mean shift was also used for *tracking of motions* and for *nonlinear spaces* through Riemannian manifolds.

Mean shift segmentation results

512
x
512



(a)



(b)

Lake image. (a) Original. (b) Segmented with $(h_s, h_r, M) = (16, 7, 40)$.

(16,7,40)

Eliminate spatial regions containing less than M pixels.

256 x 256


(8,7,20)



(a)



(b)




(c)

MIT image. (a) Original. (b) Segmented $(h_s, h_r, M) = (8, 7, 20)$. (c) Region boundaries.

The sky changes with location but is still segmented into one.

Sometimes it is possible to take planar surfaces and represent it still as a constant surface.

parameters (8,7,20)





parameters (8,7,100)

Mean shift pros and cons

- Pros
 - Does not assume spherical clusters.
 - Just two parameters (window sizes).
 - Finds variable number of modes, which are *not given*.
 - Robust to outliers and weak nonconstant regions.
- Cons
 - Output depends on window size.
 - Efficient implementation uses on short cuts in the search.
 - Does not scale well directly with dimension of feature space is above ten.